

REFERENCE

Concepts

The vocabulary that comes up in every Claude Code session — skills, CLAUDE.md, subagents, plan mode, slash commands.

Through-line: **almost all of this is just text**. Skills, commands, CLAUDE.md, even most "tools" are text files Claude reads. Get the right text in front of it and it behaves the way you want.

CLAUDE.md

A plain markdown file that acts as an **automatic, project-level system prompt**. Every time you start a session in a folder, Claude reads the CLAUDE.md at the root (and walks *up* the folder tree, reading any others along the way, up to your home directory). Use it to ground the model in *your* reality — who you are, the kind of work you produce, terms of art in your field, conventions to follow.

It's just text. One sentence ("always end with a limerick") changes behavior for an entire folder of work. Edit, **save**, and **open a new session** for changes to take effect. Generate a starter one with */init*.

Skills

A **reusable, packaged capability** you (or Anthropic) can give Claude — a folder with a SKILL.md (instructions Claude follows) plus any supporting files or scripts. Skills load *on demand*: Claude reads the short description, and when a task matches, it pulls in the full instructions. Saved expertise for recurring jobs ("draft a rubric this way," "format a handout in our house style") so you don't re-explain every time.

Subagents

A **separate Claude working in the background** on a focused task, with its own fresh context window. The main session hands off a job ("research X across these files," "review this draft"), the subagent does the work without cluttering your main conversation, and reports back just the result. Useful for big searches or parallel work — you keep the conclusion, not the mess.

Plan Mode

A working mode where Claude **researches and proposes a plan before changing anything**. Toggle it on (Shift+Tab), describe what you want, and Claude reads, thinks, and comes back with a step-by-step plan for your approval — nothing is written or run until you say go. The safest way to start any non-trivial task, and a natural fit for the idea that *a good plan is itself the deliverable*.

Slash commands

Shortcuts you type to control the session rather than to talk to the model. A command is only recognized at the **start** of a message; anything after it is passed along as arguments. Type */* in a session to see what's available.

REFERENCE

Commands worth knowing

Notation: *<arg>* = required · *[arg]* = optional. The slash-command set you'll actually reach for.

COMMAND	WHAT IT DOES
<i>/init</i>	Generate a starter CLAUDE.md so Claude understands the project. Run this first in a new folder/codebase .
<i>/clear</i>	Start a fresh context window — drop the current conversation and begin clean. (Cheap, and often the right move when a thread gets long or goes sideways.)
<i>/compact [instructions]</i>	Summarize the conversation to free up context without losing the thread. Optionally focus the summary.
<i>/context [all]</i>	Visualize where your context window is going, with warnings and tips. Pairs with <i>/compact</i> .
<i>/rewind</i>	Roll code and/or conversation back to an earlier checkpoint. (aliases: <i>/checkpoint</i> , <i>/undo</i>)
<i>/resume [session]</i>	Return to an earlier conversation by ID/name, or open the picker. (alias: <i>/continue</i>)
<i>/branch [name]</i>	Fork the current conversation , preserving the original to return to later. (alias: <i>/fork</i>)
<i>/btw <question></i>	Ask a quick side question without bloating the main conversation history.
<i>/usage</i>	Session cost, plan limits, and activity stats. (aliases: <i>/cost</i> , <i>/stats</i>)

Remember: a command is only recognized at the *start* of a message. Everything after it becomes the command's argument.

Source: Claude Code command reference — code.claude.com/docs/llms.txt