

It's all text

A model only ever reads **text**. Whatever you hand it — a PDF, a spreadsheet, a web page — is first flattened to plain text, chopped into **tokens**, turned into numbers, and read **in order**. It writes tokens back, in order. **Text in, text out.**

The diagram illustrates the process of flattening and tokenizing data for a machine learning model. It is divided into two main horizontal sections.

Top Section: Flattening and Tokenization

- 1 · ANY FORMAT:** A box containing five file format icons: .pdf, .docx, .csv, .md, and .html.
- FLATTEN:** An arrow pointing from the first box to the second box.
- 2 · PLAIN TEXT:** A box containing the text: "Shall I compare thee to a summer's day?"
- TOKENIZE:** An arrow pointing from the second box to the third box.
- 3 · TOKENS (NUMBERS):** A box containing the text "Shall I compare thee" with each word in a separate box. Below each word box is a number: "Shall" (12678), "I" (357), "compare" (9072), and "thee" (39813).


Bottom Section: Model Input and Output

- TEXT IN:** A box containing three tokens: tok, tok, tok.
- :** An arrow pointing from the "TEXT IN" box to the "THE MODEL" box.
- THE MODEL:** A box containing the text "reads left → right".
- :** An arrow pointing from the "THE MODEL" box to the "TEXT OUT" box.
- TEXT OUT:** A box containing three tokens: tok, tok, tok.

WHERE IT LIVES: TEXT FILES

STRUCTURED _____ FREE-FORM

200K – 1M tokens →



BOK CENTER · LEARNING LAB